

Ben_Kacikanis

Ben Kacikanis

15/09/2020

Information Retrieval Project

This project is an example of a basic search engine's information retrieval system. There are a list of a small corpus of documents, in addition to an example query. I will clean common meaningless words, and be using a term document matrix and take the inverse of the frequency to give words with the least frequency higher weights.

```
library(tm)
```

```
## Warning: package 'tm' was built under R version 3.6.3
```

```
## Loading required package: NLP
```

```
library(SnowballC)
```

```
## Warning: package 'SnowballC' was built under R version 3.6.3
```

```
library(slam)
```

Corpus of documents

```
doc1 <- "Stray cats are running all over the place. I see 10 a day!"
doc2 <- "Cats are killers. They kill billions of animals a year."
doc3 <- "The best food in Columbus, OH is the North Market."
doc4 <- "Brand A is the best tasting cat food around. Your cat will love it."
doc5 <- "Buy Brand C cat food for your cat. Brand C makes healthy and happy cats."
doc6 <- "The Arnold Classic came to town this weekend. It reminds us to be healthy."
doc7 <- "I have nothing to say. In summary, I have told you nothing."
```

Organizing docs and introduce query

```
doc.list<- list(doc1,doc2,doc3,doc4,doc5,doc6,doc7)
N.docs<-length(doc.list)
names(doc.list)<- paste0("doc",c(1:N.docs))
query<- "Healthy cat food"
```

```
my.docs<- VectorSource(c(doc.list,query))
my.docs$Names<- c(names(doc.list),query)
my.corpus<-Corpus(my.docs)
```

Transform corpus to remove variations of words to improve accuracy of term document matrix

```
getTransformations()
```

```
## [1] "removeNumbers"      "removePunctuation" "removeWords"
## [4] "stemDocument"       "stripWhitespace"
```

```
my.corpus<- tm_map(my.corpus,removePunctuation)
```

```
## Warning in tm_map.SimpleCorpus(my.corpus, removePunctuation): transformation
## drops documents
```

```
content(my.corpus[0])
```

```
## character(0)
```

```
my.corpus<-tm_map(my.corpus,stemDocument)
```

```
## Warning in tm_map.SimpleCorpus(my.corpus, stemDocument): transformation drops
## documents
```

```
my.corpus<-tm_map(my.corpus, content_transformer(tolower))
```

```
## Warning in tm_map.SimpleCorpus(my.corpus, content_transformer(tolower)):
## transformation drops documents
```

```
my.corpus<-tm_map(my.corpus,stripWhitespace)
```

```
## Warning in tm_map.SimpleCorpus(my.corpus, stripWhitespace): transformation drops
## documents
```

Term Document Matrix

```
term.document.matrix.stm<- TermDocumentMatrix(my.corpus)
colnames(term.document.matrix.stm)<-c(names(doc.list),"query")
inspect (term.document.matrix.stm)
```

```
## <<TermDocumentMatrix (terms: 46, documents: 8)>>
## Non-/sparse entries: 62/306
## Sparsity           : 83%
## Maximal term length: 8
## Weighting          : term frequency (tf)
## Sample            :
##           Docs
## Terms      doc1 doc2 doc3 doc4 doc5 doc6 doc7 query
```

```
## are      1      1      0      0      0      0      0      0
## best     0      0      1      1      0      0      0      0
## brand    0      0      0      1      2      0      0      0
## cat      1      1      0      2      3      0      0      1
## food     0      0      1      1      1      0      0      1
## have     0      0      0      0      0      0      2      0
## healthi  0      0      0      0      1      1      0      1
## noth     0      0      0      0      0      0      2      0
## the      1      0      2      1      0      1      0      0
## your     0      0      0      1      1      0      0      0
```

```
term.document.matrix<- as.matrix(term.document.matrix.stm)
cat("Dense matrix representation costs", object.size(term.document.matrix), "bytes.\n",
    "Simple triplet matrix representation costs", object.size(term.document.matrix.stm),
    "bytes.")
```

```
## Dense matrix representation costs 7184 bytes.
## Simple triplet matrix representation costs 6440 bytes.
```

tfidf.matrix, inverse matrix reflects importance of words with less frequency and higher meaning

```
get.tf.idf.weights<-function(tf.vec){
  n.docs<- length(tf.vec)
  doc.frequency<-length(tf.vec[tf.vec>0])
  weights<- rep(0,n.docs)
  weights[tf.vec>0]<-(1+log2(tf.vec[tf.vec>0]))*(1+log2(n.docs/doc.frequency))
  return (weights)
}
```

```
get.tf.idf.weights(c(1, 2, 3, 0, 0, 6))
```

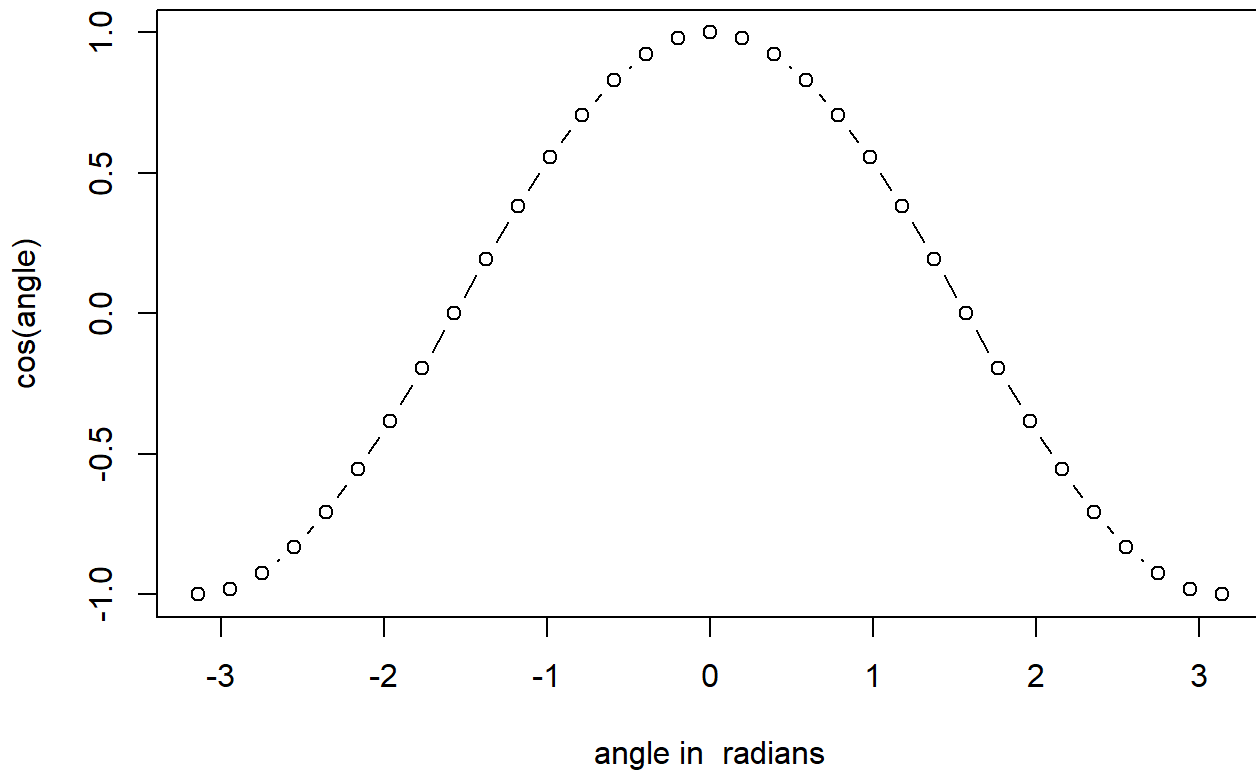
```
## [1] 1.584963 3.169925 4.097069 0.000000 0.000000 5.682031
```

```
tfidf.matrix<- t(apply(term.document.matrix,1, FUN= function(row){get.tf.idf.weights(row)}))
colnames(tfidf.matrix)<- colnames(term.document.matrix)
```

Cosine similarity to give measure how similar documents are

```
angle<-seq(-pi,pi,by = pi/16)
plot(cos(angle)~angle,type="b", xlab = "
angle in radians", main ="Cosine similarity by
angle")
```

Cosine similarity by angle



```
tfidf.matrix<- scale(tfidf.matrix,center = FALSE,  
                     scale = sqrt(colSums(tfidf.matrix^2)))  
tfidf.matrix[0:3,]
```

```
##  
## Terms      doc1      doc2 doc3      doc4      doc5 doc6 doc7      query  
## all 0.3538079 0.0000000 0 0.0000000 0.0000000 0 0 0.0000000  
## are 0.2653559 0.2889215 0 0.0000000 0.0000000 0 0 0.0000000  
## cat 0.1484288 0.1616103 0 0.3196129 0.3499457 0 0 0.4718391
```

```
query_vector<- tfidf.matrix[,N.docs+1]  
tfidf.matrix<- tfidf.matrix[,1:N.docs]
```

Results of how similar query is to documents.

```
doc_scores<- t(query_vector) %*% tfidf.matrix  
  
results.df<- data.frame(doc=names(doc.list),score= t(doc_scores),text=unlist(doc.list))  
results.df<- results.df[order(results.df$score,decreasing = T),]
```

```
options(width = 200)  
print(results.df,row.names=F,right=F,digits=2)
```

```
## doc score text  
## doc5 0.388 Buy Brand C cat food for your cat. Brand C makes healthy and happy cats.  
## doc4 0.258 Brand A is the best tasting cat food around. Your cat will love it.
```

```
## doc6 0.149 The Arnold Classic came to town this weekend. It reminds us to be healthy.  
## doc3 0.128 The best food in Columbus, OH is the North Market.  
## doc2 0.076 Cats are killers. They kill billions of animals a year.  
## doc1 0.070 Stray cats are running all over the place. I see 10 a day!  
## doc7 0.000 I have nothing to say. In summary, I have told you nothing.
```